





MKDA density, NiMARE default kernel, FDR-independent correction; maps regenerated identically for every arm.

R^2 is over voxels nonzero in either map ($r2_nonzero$), so it measures agreement about findings rather than about empty space. The repo's convention correlates every finite voxel, which on these sparse maps is 88-96% voxels zero in both, and averages 0.533 against 0.236 here; a brain mask removes almost none of that, because the zeros are inside the brain. All three are in `data/<project>_maps.csv`.

Panel b averages each project's mapped analyses, counted under its label, and offsets the arms horizontally so coincident points stay visible. No mapped analysis comes from a paper its own arm rejected. Four of the five projects annotate from each arm's own document; VBM substance use annotates from title and abstract only, because its full-text arm's articles are not on this host.

Re-running one project's annotation unchanged moved its arms by 0.034 R^2 , which exceeds every between-arm gap shown in b.

The two query arms are not a model: the review's own inclusion criteria run as predicates over the extraction record, then its own contrast over that paper's analyses. Strict drops what a record cannot answer and permissive keeps it. Emotion regulation has no query arm, and maps built from fewer than 5 analyses are not scored.